Inputs

16 frames per video, sampled evenly

Speech transcript, faster-whisper large-v3-turbo

Question; for Task 2 also its released options

Task 1, answer generation

Sample S times
S = 8 or 16; each sees a random 8 of the 16 frames

Qwen3-VL-8B
asked for the single best short answer

Cluster, rank
MiniLM clusters; rank by how often an answer recurs

Slate of ten
distinct answers, best first

cons8b_s8, cons8b_s16: S = 8, 16. fewshot8b: S = 8 + the four nearest 2025 examples.
all8b: + a transcript for verbal questions. v3control: one call for all ten answers.

Task 2, multiple choice

Plan
8, 12 or 16 frames; transcript if the question is audio

Three votes
each ranks all options (T = 0, 0.7 and 0.7)

Combine
MRR-weighted Borda over votes; re-ask unparsable

Ranked options
a permutation of the released options

mc32b, mc32bdense: Qwen3-VL-32B. mc30b: Qwen3-VL-30B-A3B. mc122b: Qwen3.5-122B.
mcfusion: reciprocal-rank fusion of the 32B, 122B and 30B rankings.
Withheld: a two-tier cascade that escalates likely 8B errors to 32B.